Mastering Amazon SageMaker

MLOps

Name

Role

Email

Agenda

- Amazon Sagemaker MLOps Capabilities
  - SageMaker Pipelines
  - SageMaker Model Registry
  - SageMaker Projects
- Demo
- Q&A

aws machine learning

Amazon SageMaker MLOps capabilities

aws machine learning

Operational challenges with managing the ML lifecycle

Manual iterative processes slow down ML innovation

Difficult to scale the number of models in production

CI/CD for ML requires writing custom code

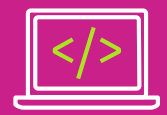

Compliance requirements are difficult to meet

aws machine learning

Amazon SageMaker MLOps

Streamline the ML lifecycle

Automate ML workflows to scale model development

Build CI/CD pipelines for ML to accelerate model deployment

Catalog model versions, metadata, metrics, and approvals for traceability and reusability

Track lineage for troubleshooting and compliance

Maintain accuracy of predictions after models are deployed

Enhance governance and security

aws machine learning

Amazon SageMaker MLOps

Components

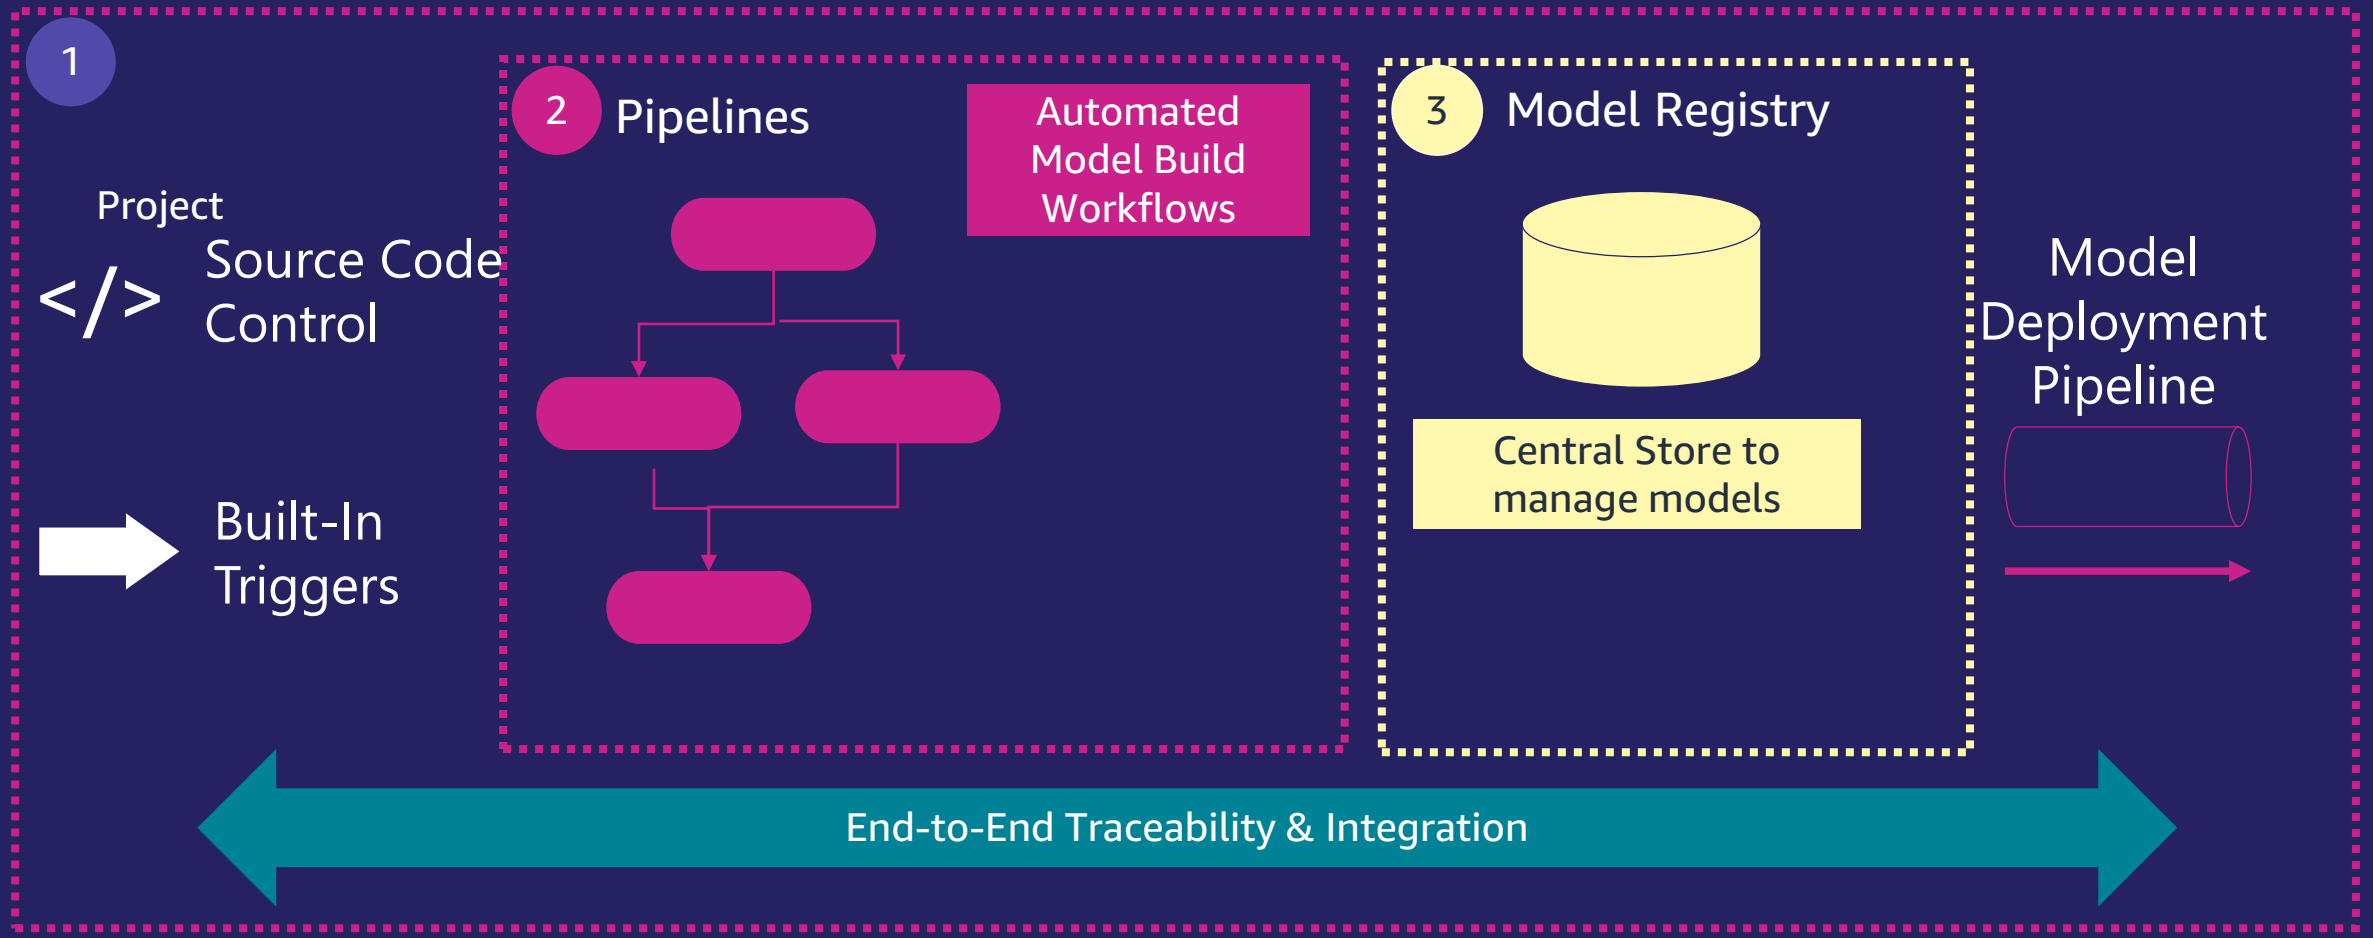

aws machine learning

Amazon SageMaker Pipelines

Components

- **Flexible:**
  - ✓ Use Projects for end-to-end pipelines that incorporate CI/CD practices
  - ✓ Optionally, use Pipelines or Model Registry without Projects to meet the needs of your use case

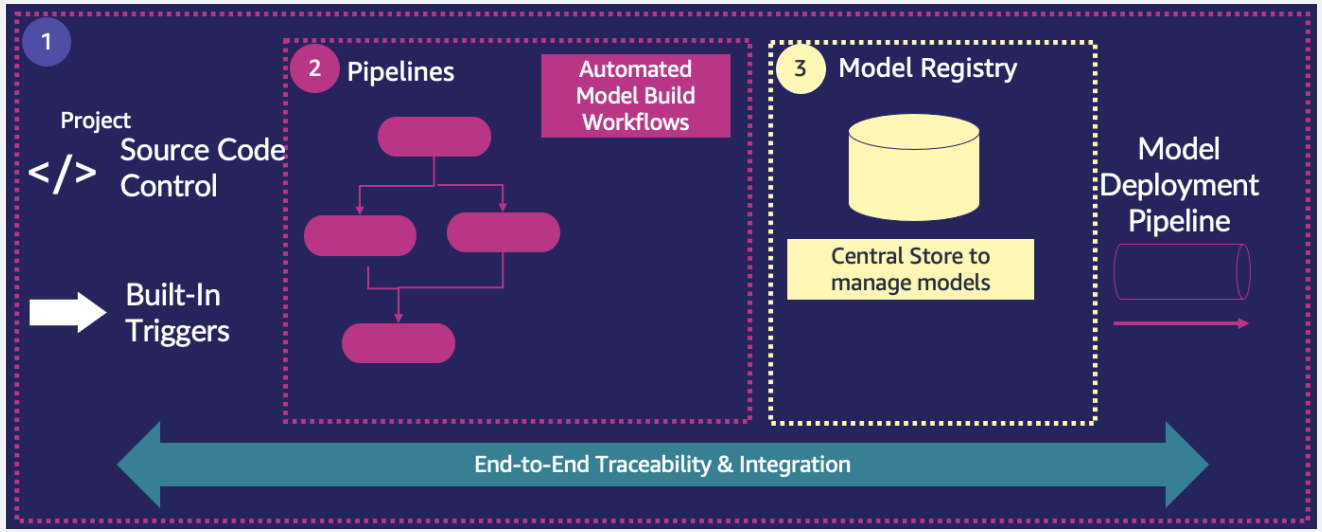

Amazon SageMaker Pipelines

Components

Extensible:

- ✓ Take advantage of built-in MLOps Project Templates

or

- ✓ Create custom MLOps Project Templates to meet the unique demands of your corporate or regulatory requirements

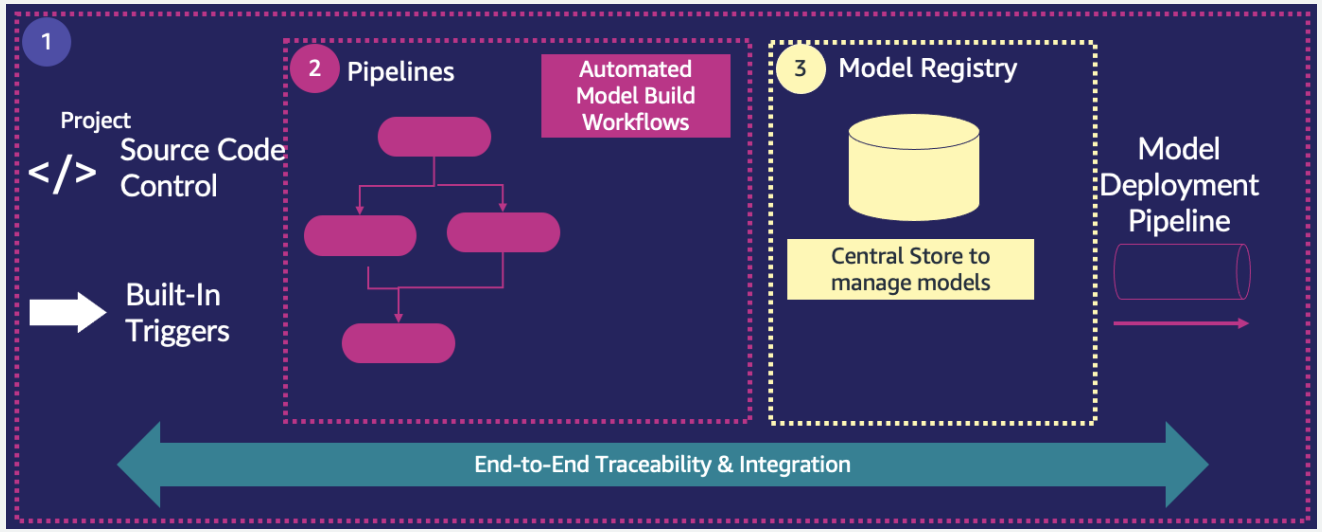

aws machine learning

SageMaker Pipelines

Creating Amazon SageMaker Automated Pipelines

Pipeline-as-Code

How it Works …

1 Create Steps →

Define & configure each step

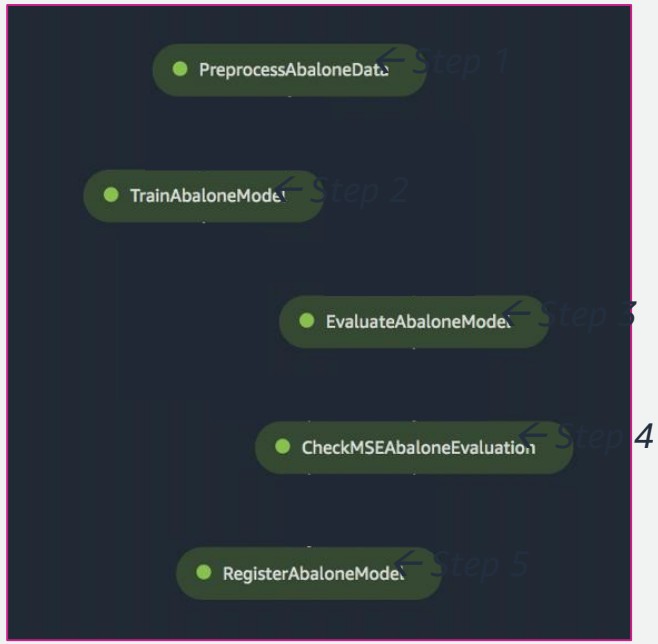

2 Define Pipeline →

Define & configure the workflow

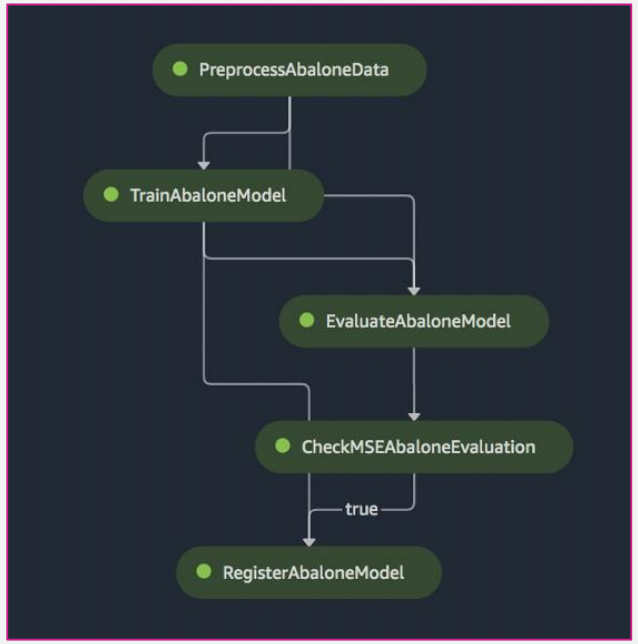

3 Start Pipeline →

Execute the pipeline

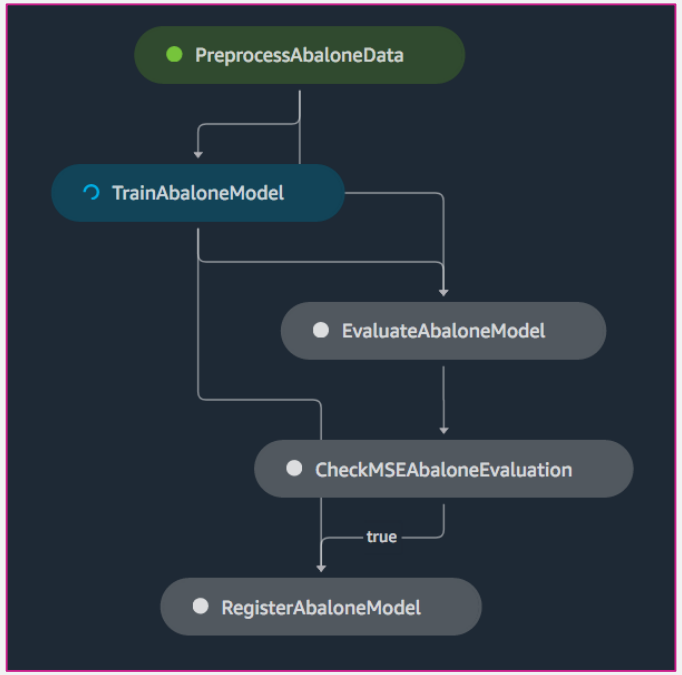

aws machine learning

Amazon SageMaker Pipelines

Support for Step Caching →

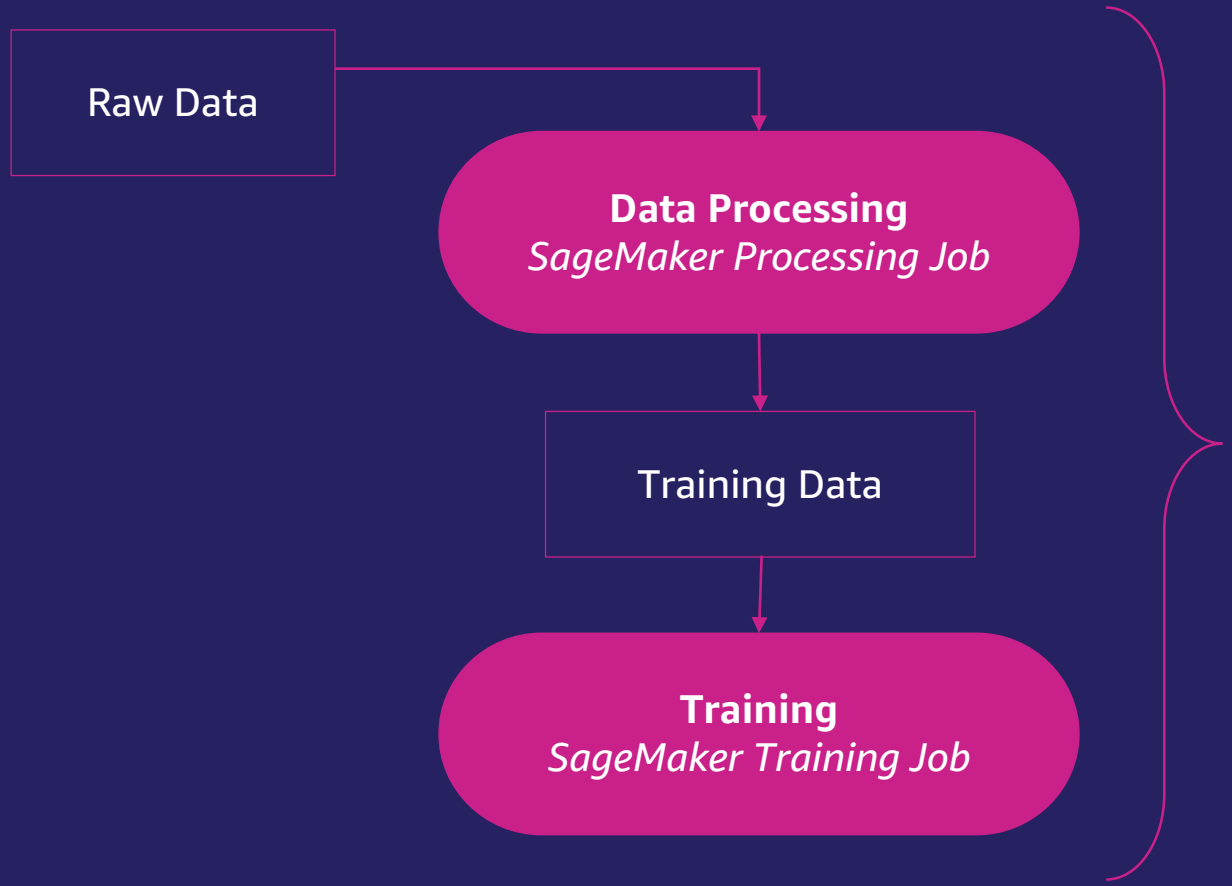

What if you want to:
- tweak hyperparameters?
- modify training code?

aws machine learning

Amazon SageMaker Pipelines

Support for Pipeline Parameters

1 Configure your parameter →

```
from sagemaker.workflow.parameters import (
    ParameterInteger,
    ParameterString,
    ParameterFloat
)

processing_instance_count = ParameterInteger(
    name="ProcessingInstanceCount",
    default_value=1
)
```

2 Pass in parameter on pipeline create →

```
pipeline = Pipeline(
  name=pipeline_name,
  parameters=[
    processing_instance_count
  ],
steps=[step_process]
)
```

3 Optionally, pass a non-default value for pipeline execution →

```
execution = pipeline.start(
   parameters=dict(
            ProcessingInstanceType="ml.c5.xlarge
",          ModelApprovalStatus="Approved”
   )
)
```

aws machine learning

SageMaker Model Registry

aws machine learning

Amazon SageMaker Pipelines

Model Registry

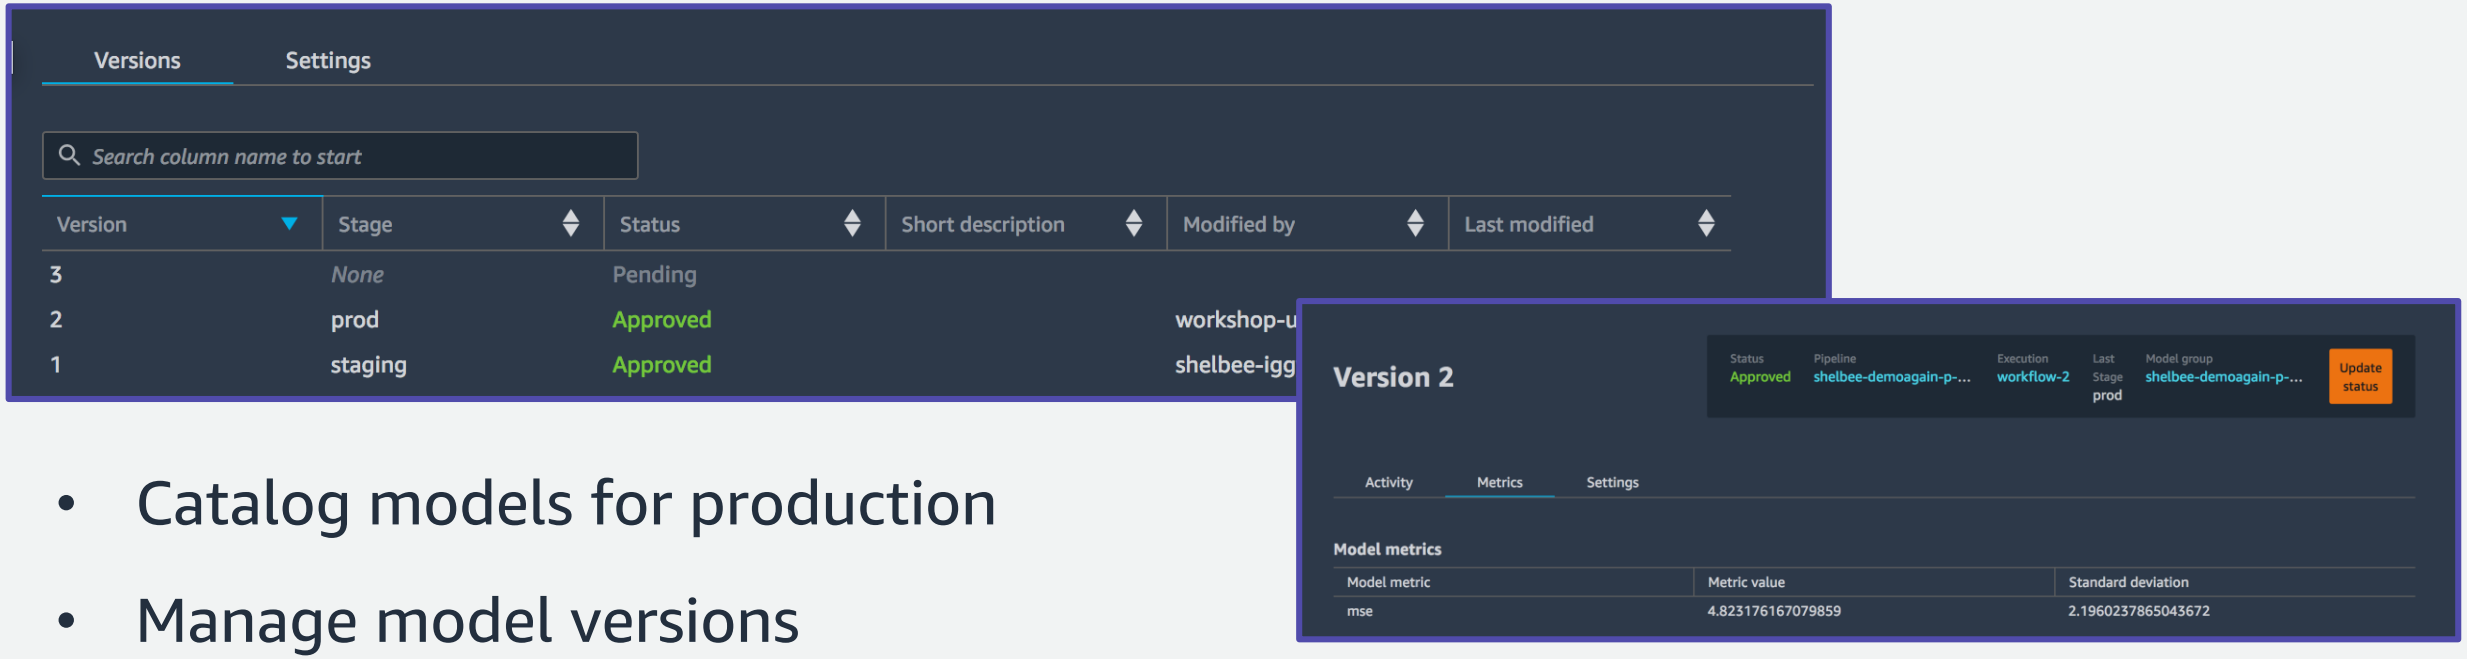

- Catalog models for production
- Manage model versions
- Associate metadata with a model
- Manage the approval status of a model
- Deploy models to production (with Projects)
- Track model performance metrics
- Rollback model versions

aws machine learning

Amazon SageMaker Pipelines

Model Registry

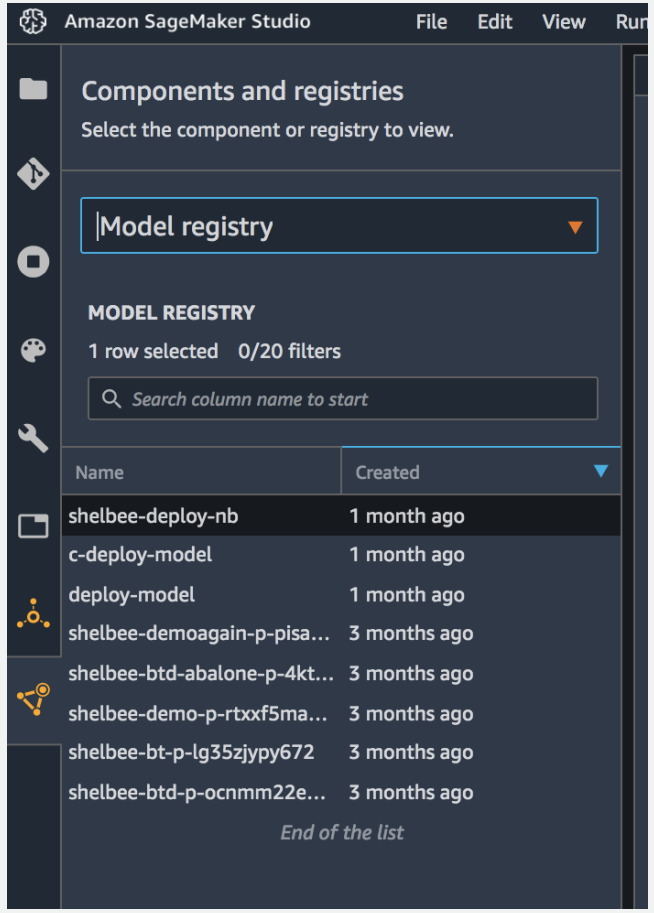

Hierarchy →

Model Registry

Model Group (1)	Model Group (2)	Model Group (n)
vn	vn	v1
v1	v2	
	v1	

aws machine learning

Amazon SageMaker Pipelines

Model Registry

How it Works …

1 **Create Model Group →**

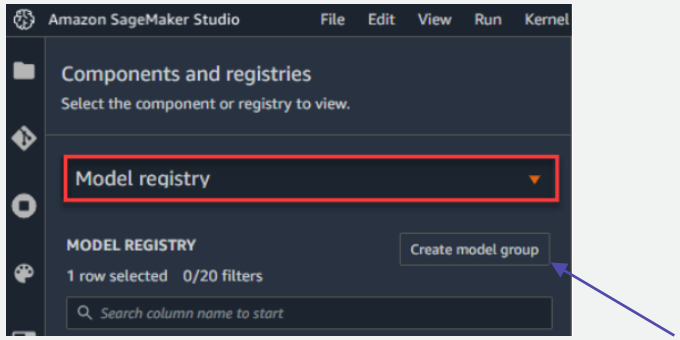

2 **Register a Model Version →**

Within a Pipeline, using RegisterModel Step:

```
step_register = RegisterModel(
    name="RegisterAbaloneModel",
    estimator=xgb_train,
    model_data=step_train.properties.ModelArtifacts.S3ModelArtifacts,
    content_types=["text/csv"],
    response_types=["text/csv"],
    inference_instances=["ml.t2.medium", "ml.m5.large"],
    transform_instances=["ml.m5.large"],
    model_package_group_name=model_package_group_name,
    approval_status=model_approval_status,
    model_metrics=model_metrics,
)
```

=

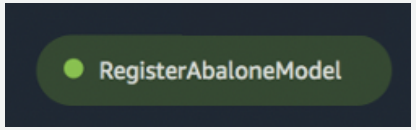

~OR~

Using boto3:

```
create_model_package_response = sm_client.create_model_package(**create_model_package_input_dict)
model_package_arn = create_mode_package_response["ModelPackageArn"]
print('ModelPackage Version ARN : {}'.format(model_package_arn))
```

aws machine learning

SageMaker Projects

aws machine learning

Amazon SageMaker Pipelines - Projects
Getting Started

One-Time Setup

1 *Enable Projects for your Studio Domain*

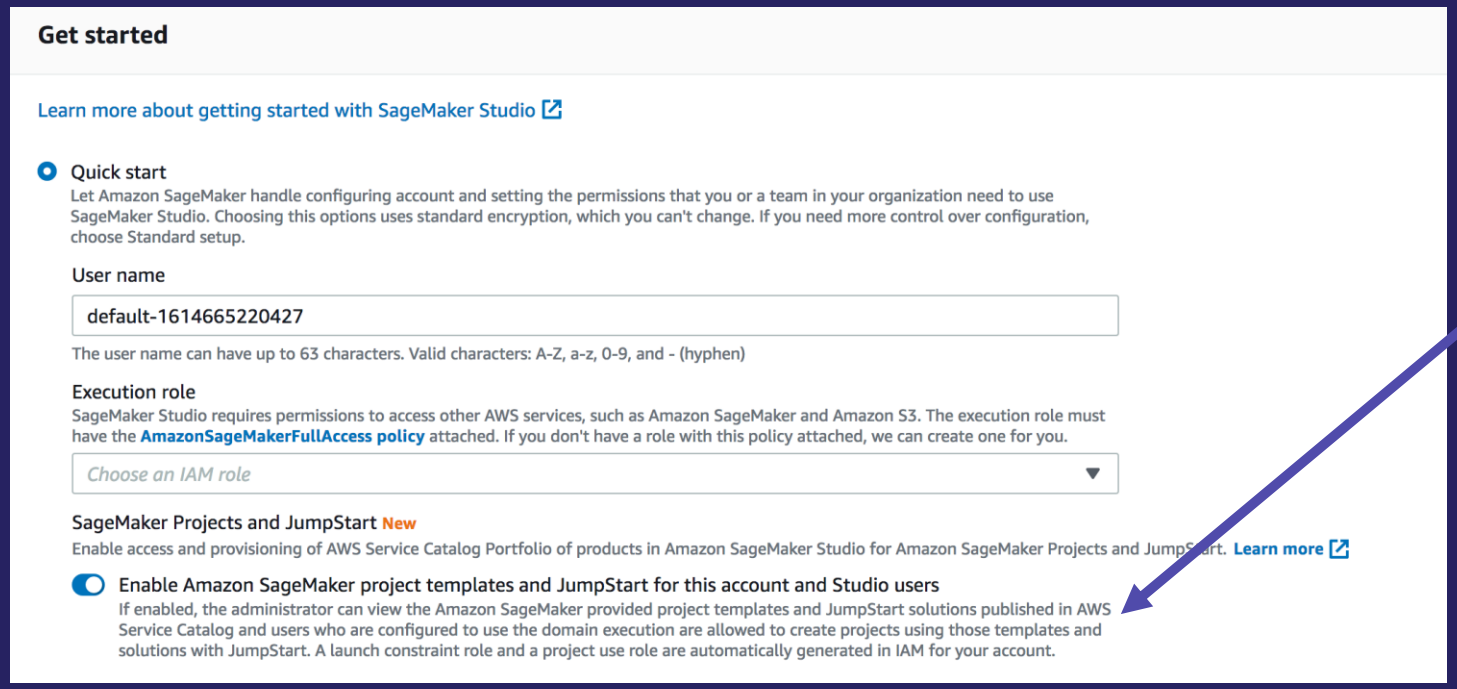

Enable on a new domain or update an existing domain

aws machine learning

Amazon SageMaker - Projects
Features

- Utilize Built-In MLOps Project Templates:
  1. Build, Train, Deploy
  2. Build, Train
  3. Deploy

- Create Custom MLOps Project Templates

aws machine learning

Amazon SageMaker Pipelines - Projects

Built-In ML-Ops Project Templates

How it Works …

1 *Select Create Project*

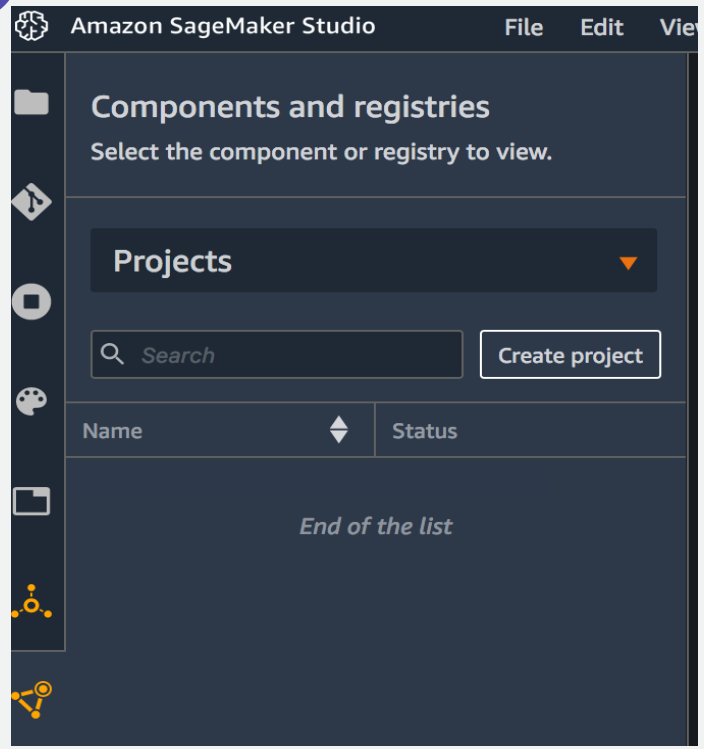

2 *Select Template*

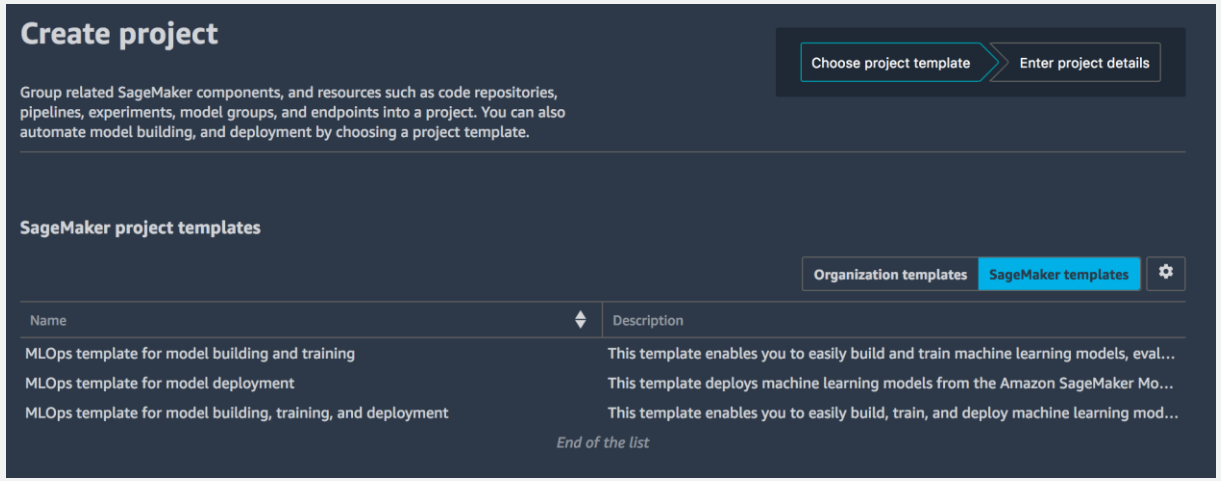

Everything needed to create a CI/CD Pipeline for Machine Learning gets automatically provisioned & configured for you……

aws machine learning

Amazon SageMaker Pipelines - Projects

High Level Services View

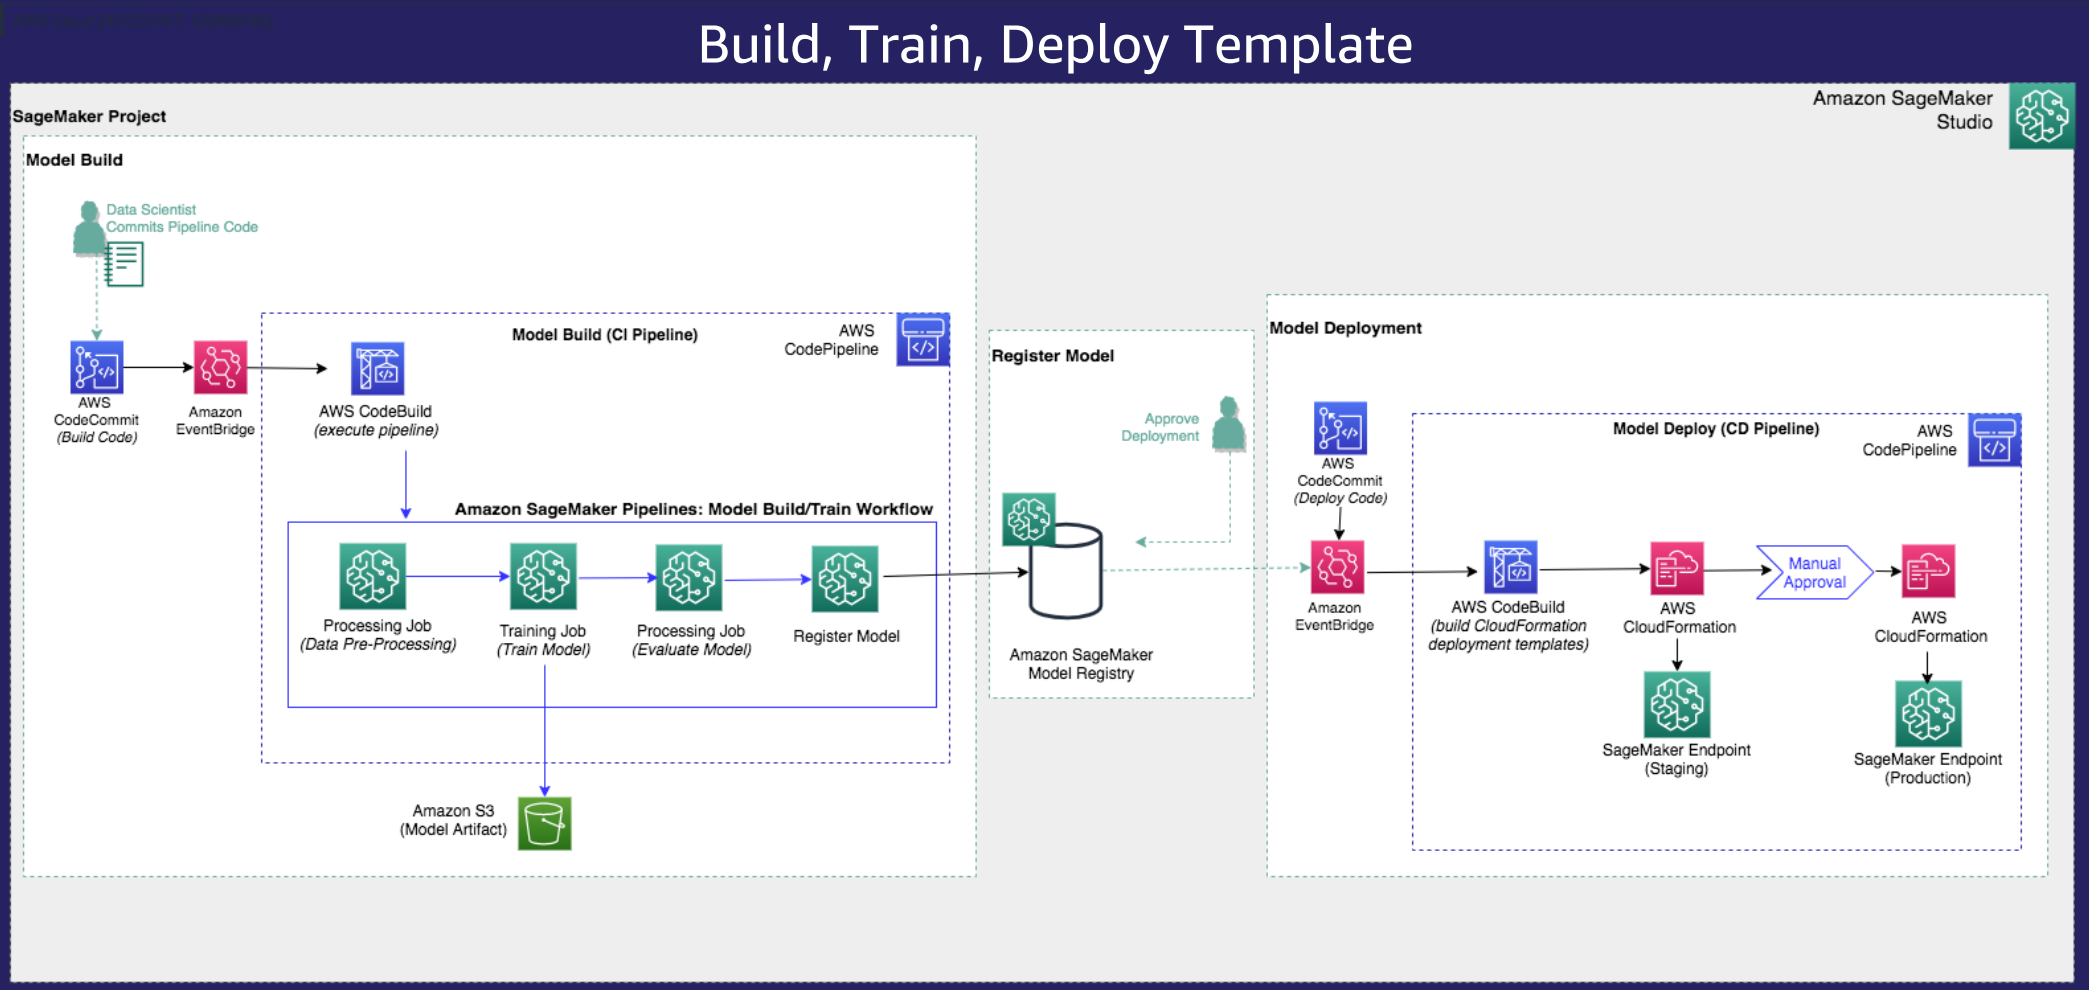

Amazon SageMaker Pipelines - Projects

Underlying Services

AWS Service Catalog

Description: *AWS Service Catalog allows organizations to create and manage catalogs of IT services that are approved for use on AWS.*

- Built-In Project templates are offered through Products contained in a managed AWS Service Catalog Portfolio

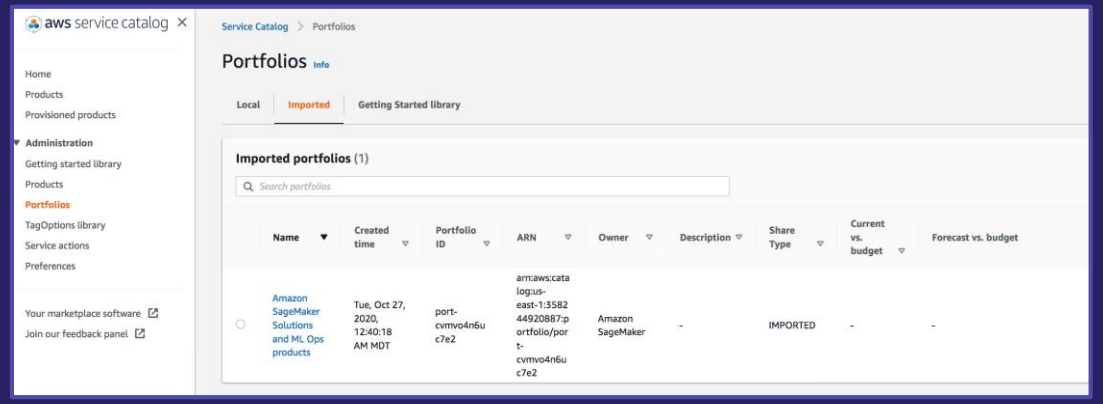

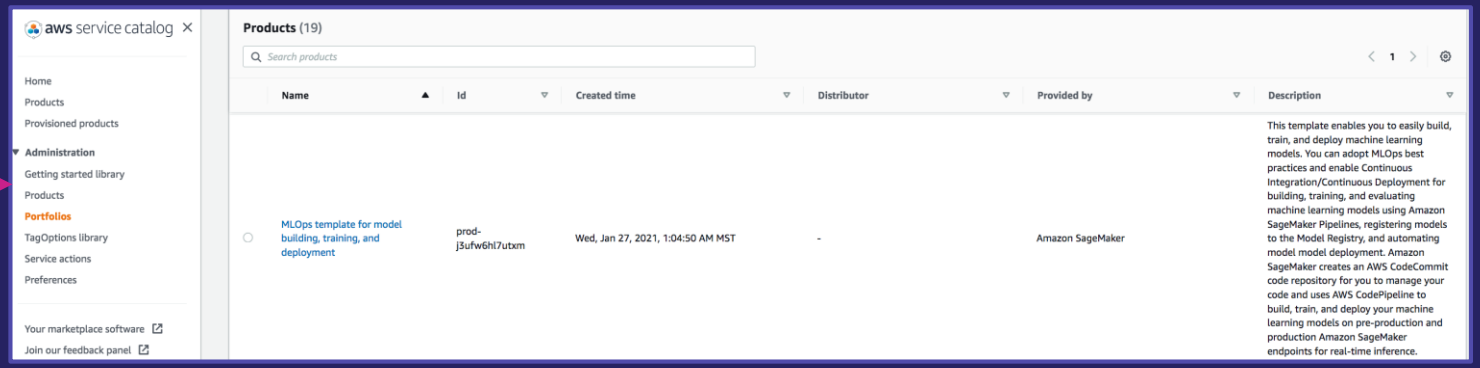

aws machine learning

Amazon SageMaker Pipelines - Projects
Underlying Services

AWS CloudFormation

Description: *AWS CloudFormation templates allow for the ability to consistently provision, configure, and manage resources using Infrastructure-as-Code*

- SageMaker Pipelines uses AWS CloudFormation to provision & configure:

1. Project resources needed to build an end-to-end pipeline

2. Model endpoint deployments for an instantiated Project

aws machine learning

Amazon SageMaker Pipelines - Projects
Underlying Services

AWS CloudFormation

1. Project resources needed to build an end-to-end pipeline

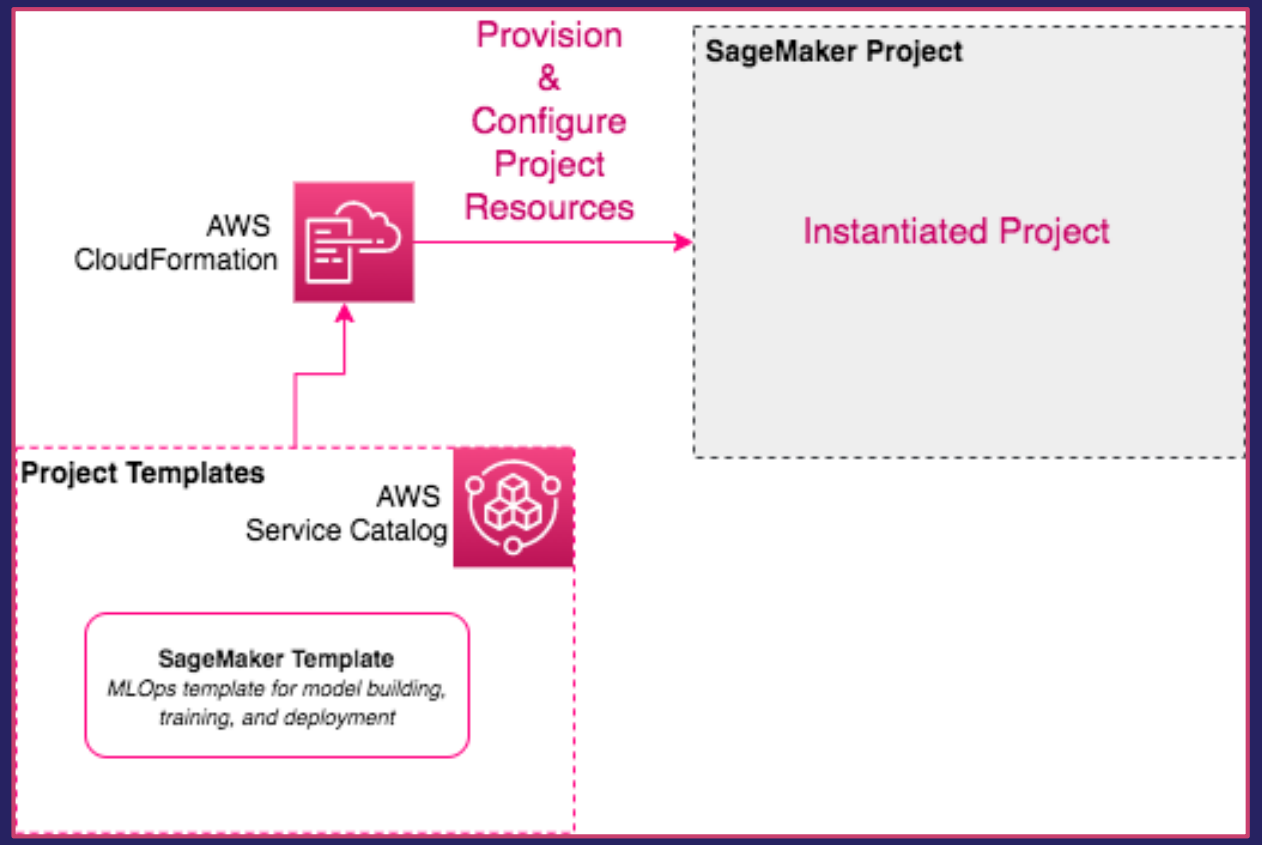

aws machine learning

Amazon SageMaker Pipelines - Projects

Underlying Services

AWS CodePipeline

Description: *AWS CodePipeline is a fully managed Continuous Delivery (CD) service*

- SageMaker Pipelines uses AWS CodePipeline to orchestrate model build and model deploy activities

1. Model Build:

2. Model endpoint deployments for an instantiated Project

aws machine learning

Amazon SageMaker Pipelines - Projects Features

- Utilize Built-In MLOps Project Templates:
  1. Build, Train, Deploy
  2. Build, Train
  3. Deploy

- Create Custom MLOps Project Templates

aws machine learning

Amazon SageMaker Pipelines - Projects
Building Custom Templates

1 *Create Your Custom Template*

AWS CloudFormation

Required parameters to include:

```
SageMakerProjectName:
    Type: String
    Description: Name of the project

  SageMakerProjectId:
    Type: String
    Description: Service generated Id of the project.
```

Tip: Use the SageMaker built-in templates as a starting point

2 *Upload CloudFormation Template to S3*

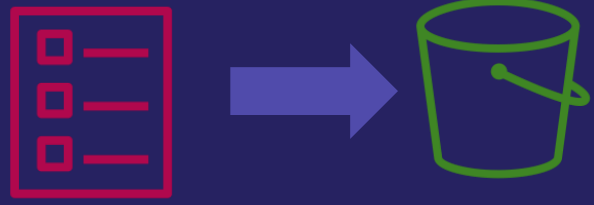

Amazon S3

aws machine learning

Amazon SageMaker Pipelines - Projects
Building Custom Templates

Service Catalog Tasks

AWS Service Catalog

3 *Create Portfolio*

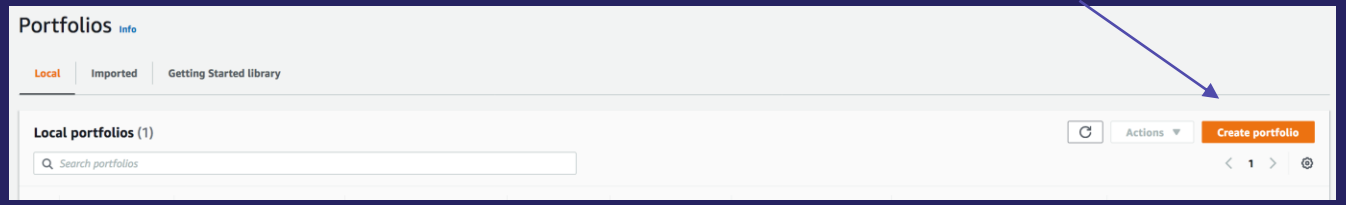

4 *Create Product*

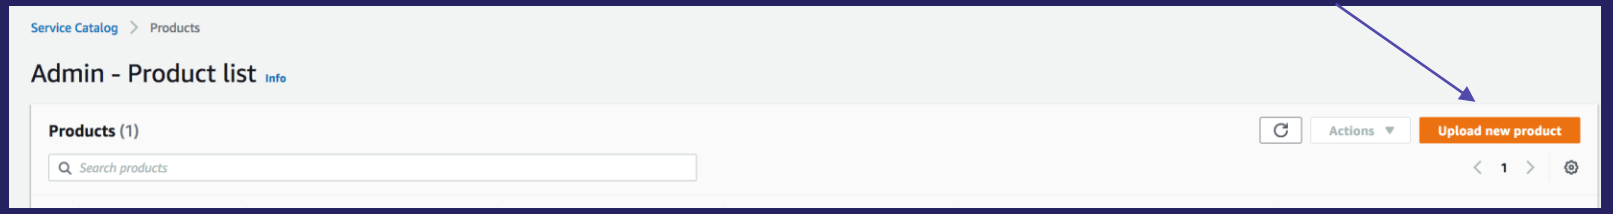

Required Tags

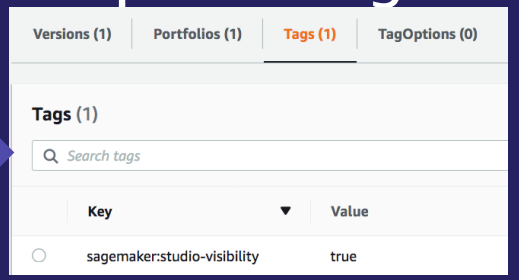

Your template will now be available for use in SageMaker Studio!

aws machine learning

Amazon SageMaker

Next Steps

Onboarding & Processing

- https://docs.aws.amazon.com/sagemaker/latest/dg/gs-studio-onboard.html
- https://docs.aws.amazon.com/sagemaker/latest/dg/processing-job.html

Training

- https://docs.aws.amazon.com/sagemaker/latest/dg/train-model.html
- https://docs.aws.amazon.com/sagemaker/latest/dg/distributed-training.html
- https://aws.amazon.com/sagemaker/debugger

Deployment

- https://docs.aws.amazon.com/sagemaker/latest/dg/realtime-endpoints.html
- https://docs.aws.amazon.com/sagemaker/latest/dg/serverless-endpoints.html
- https://docs.aws.amazon.com/sagemaker/latest/dg/async-inference.html
- https://docs.aws.amazon.com/sagemaker/latest/dg/batch-transform.html

https://github.com/aws/amazon-sagemaker-examples

https://sagemaker.readthedocs.io/en/stable/index.html

aws machine learning

Q & A

Please Complete the session Survey

Thank you!